ML - Capstone Project

- Project Phases: Data Exploration and Understanding: Dive into the dataset to understand the landscape of laptop specifications. Visualize trends in laptop prices and identify potential influential features. Data

- Preprocessing: Handle missing values, outliers, and encode categorical variables. Ensure the dataset is ready for model training.

- Feature Engineering: Extract meaningful features to enhance model performance. Consider creating new features that capture the essence of laptop pricing. Model Development: Employ machine learning algorithms such as Linear Regression, Random Forest, and Gradient Boosting to predict laptop prices. Evaluate and choose the model that aligns best with the project's objectives.

- Hyperparameter Tuning: Fine-tune the selected model to achieve optimal performance. Real-time Predictions: Implement a mechanism for the model to make predictions for new laptops entering the market.

- Interpretability and Insights: Uncover insights into which features play a pivotal role in pricing decisions. Ensure that SmartTech Co. can interpret and trust the model's predictions.

1. **Introduction**	**Objective**: Develop a machine learning model to predict laptop prices based on various features.	**Scope**: Provide insights into factors influencing laptop prices and support SmartTech Co. in market positioning and strategy.	2. **Data Exploration and Understanding**	**Dataset Overview**: Describe the dataset, including key features and their types (e.g., categorical, numerical).	**Exploratory Data Analysis (EDA)**: Share visualizations and statistics that highlight patterns, correlations, and distributions in the data.	3. **Data Preprocessing**
Handling Missing Values: Explain how missing values were addressed.	**Feature Encoding**: Detail the one-hot encoding process for categorical features.	**Feature Scaling**: Discuss any scaling techniques applied.	4. **Model Development**	**Algorithms Used**: • **Linear Regression**: Explain its role and how it was implemented. • **XGBoost**: Describe its usage and advantages over other algorithms.	**Model Training**: Discuss how the models were trained, including hyperparameter tuning.	5. **Model Performance**
Evaluation Metrics: Present metrics such as Mean Absolute Error (MAE), Mean Squared Error (MSE), R-squared, and any other relevant metrics.	**Comparison of Models**: Compare the performance of Linear Regression and XGBoost.	6. **Insights**	**Feature Importance**: Identify which features have the most significant impact on laptop prices.	**Brand Influence**: Analyze whether the brand significantly influences the price.	**Specification Impact**: Compare how well the model performs for high-end vs. budget laptops.	**Performance on Lesser-Known Brands**: Evaluate the model's accuracy for laptops from lesser-known brands.
New Releases: Discuss the model's performance on newly released laptops not in the training dataset.	7. **Challenges and Limitations**	**Data Limitations**: Address any gaps or limitations in the dataset.	**Model Limitations**: Discuss any challenges faced during model development and limitations in predictions.	8. **Feedback and Improvements**	**Gather Feedback**: Ask stakeholders for their insights and concerns about the model's performance and applicability.	**Future Improvements**: Propose potential improvements based on feedback and observed limitations.

Feature Encoding

```
from sklearn.preprocessing import OneHotEncoder
import pandas as pd
# Example for one-hot encoding categorical features
encoder = OneHotEncoder(drop='first')
encoded_features = encoder.fit_transform(data[['Brand', 'Category']]).toarray()
encoded_df = pd.DataFrame(encoded_features, columns=encoder.get_feature_names_out())
data = data.join(encoded_df).drop(['Brand', 'Category'], axis=1)
```

Model Training

- from sklearn.linear_model import LinearRegression
- from xgboost import XGBRegressor
- from sklearn.model_selection import train_test_split
- from sklearn.metrics import mean_squared_error, r2_score

Split data

- X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

- # Train Linear Regression model
- lr_model = LinearRegression()
- lr_model.fit(X_train, y_train)

Train XGBoost model

- xgb_model = XGBRegressor()
- xgb_model.fit(X_train, y_train)

Predictions

- lr_predictions = lr_model.predict(X_test)
- xgb_predictions = xgb_model.predict(X_test)

Evaluation

- lr_mse = mean_squared_error(y_test, lr_predictions)
- xgb_mse = mean_squared_error(y_test, xgb_predictions)
- lr_r2 = r2_score(y_test, lr_predictions)
- xgb_r2 = r2_score(y_test, xgb_predictions)

- print(f'Linear Regression MSE: {lr_mse}, R2: {lr_r2}')
- print(f'XGBoost MSE: {xgb_mse}, R2: {xgb_r2}')

Final Presentation Tips

- **Visualizations**: Use charts and graphs to make your data and results more comprehensible.
- **Clarity**: Ensure that your explanations are clear and tailored to the audience's level of technical expertise.
- **Interaction**: Encourage questions and discussions to engage stakeholders and gather valuable feedback.
- This structure should help you deliver a comprehensive and insightful presentation to SmartTech Co.